Cheaper AI. What about the machines?

Why I revisited agent sandboxes.

MODEL	MACHINE	RESULT
instructions	execution	verified output

Cost, privacy, features and customization.

New sandboxes. Architecture rethink.

GOOGLE / 9 SEPTEMBER 2026

Shell + computer-use sandboxes reach GA.
Pause idle compute while preserving files.

Google release

DATABRICKS / 16 JUNE 2026

Sandbox announced within Agent Bricks.
Dedicated offering currently no-charge Beta.

Databricks news

These developments prompted my side-project architecture review, mainly around cost.
New features aren't benchmark results.

What survives?

Files, durable disk and memory are separate promises.

01

BETWEEN CALLS

Can the next call read the files?

02

AFTER STOP

Does disk remain? For how long?

03

ON RESUME

Are variables and processes restored?

E2B documents memory + files on default pause.
Our E2B/Modal checks verified files across calls.

Stop/resume was not tested in this pilot.

E2B persistence

Bring your tools. Choose the hardware.

IMAGE	ALLOCATION
Libraries and toolchain	**CPU, RAM and device**

E2B Templates; CPU-only sandboxes.

MODAL GPU offered; plan for interruptions.

DAYTONA Shared GPU placement ignores preference.

Check the exact runtime class and controls.
GPU configurations were not benchmarked here.

E2B GPU FAQ **Modal resources** **Daytona regions**

An EU region is one part of the map.

MODAL EXAMPLE / DOCUMENTED DATA LOCATIONS

COMPUTE
EEA-pinned

EEA is broader than the EU. Pinning covers compute.

LOGS
US stdout / stderr

SNAPSHOTS
US, when created

MANAGED STORAGE
US durable products

Also trace backups, control-plane data, support access and model calls. A model endpoint has its own processing boundary.

Modal residency

The machines have their own economics.

CHEAPER ACCESS TO CAPABILITY

280-fold+

GPT-3.5-level inference cost drop.
Nov 2022 to Oct 2024 (Stanford AI Index).

2026 INFRASTRUCTURE PRICE PRESSURE

Server memory supply remains tight.
TrendForce forecasts further price rises.
Hetzner raised new-order cloud prices in June.

These are different markets and time periods.

Stanford AI Index	Memory forecast	Hetzner prices

Cheaper tokens. Still a machine bill.

Same model budget. Different billed running time. Hypothetical USD usage, before credits.

MODEL-AGNOSTIC SCENARIO / NOT A MODEL RATE

Assumed model budget: $0.003 in both

E2B: 2 VCPU + 4 GIB / BILLED RUNNING TIME

10 seconds	**$0.00046**
5 minutes	**$0.01380**

Either bill can dominate. Count the whole job. Sandbox storage included; paid fees, tax, retries and other service charges excluded.

Replace with your model bill **E2B pricing**

Five questions worth asking.

01 What survives between calls, after a stop and on resume? Files, disk, memory.

02 Does a custom image also reserve a GPU? It doesn't. Request hardware separately.

03 Where do logs, backups and model calls live? Map every data surface and access path.

04 What happens when a request times out? Verify cancellation and recovery.

05 What's the cost per verified completion? Keep credits apart from gross usage.

Gross campaign cost / verified completed tasks

A benchmark for my next decision.

File/data checks produced the expected results with E2B, Modal and Google's default code. The tested setups give me a starting shortlist.

THE DEEP DIVE

Article: architecture, costs and provider trade-offs.

Repository: exact setups, workload checks and retained results.

Theo: cheaper capability can enable more work. Article + repository links in the post.

Theo's video